

# Technical Architecture Career Learning Guide

## Learning Path Overview

This guide is structured in 4 progressive phases, each building upon the previous one. The estimated timeline is 12-18 months of dedicated study and practice.

---

## Phase 1: Foundation (Months 1-3)

*Build core understanding of systems and data*

### 1.1 Database & Storage Fundamentals

#### Week 1-2: Core Database Concepts

- **SQL vs NoSQL:** Learn when to use relational vs document/key-value stores
- **ACID Properties:** Atomicity, Consistency, Isolation, Durability
- **BASE Model:** Basically Available, Soft state, Eventually consistent
- **OLAP vs OLTP:** Analytical vs transactional processing patterns

#### Week 3-4: Database Internals

- **Indexing & B+ Trees:** How databases optimize queries
- **Transaction Isolation:** Read uncommitted, committed, repeatable read, serializable
- **WAL (Write Ahead Log):** Understanding database durability mechanisms
- **Read/Write Patterns:** Optimizing for different access patterns

#### Week 5-6: Data Management

- **Data Modeling:** Normalization, denormalization strategies
- **Data Partitioning:** Horizontal and vertical partitioning
- **Data Retention:** Archival strategies and compliance
- **Hot/Cold Storage:** Cost-effective data lifecycle management

#### Practical Exercises:

- Set up PostgreSQL and MongoDB instances
- Design schemas for e-commerce and social media use cases
- Implement different isolation levels and observe behavior
- Practice query optimization with EXPLAIN plans

### 1.2 Storage Systems

#### Week 7-8: Modern Storage

- **Object Storage & S3 Basics:** Understanding cloud storage paradigms
- **File Systems:** Traditional vs distributed file systems
- **Backup & Restore:** Strategies for data protection
- **Bloom Filters:** Probabilistic data structures for efficiency

#### **Practical Exercises:**

- Deploy MinIO (S3-compatible storage)
  - Implement backup strategies with different RTO/RPO requirements
  - Build a simple bloom filter implementation
- 

## **Phase 2: Distributed Systems Core (Months 4-7)**

### **2.1 Scaling Fundamentals**

#### **Week 9-10: Basic Scaling Concepts**

- **Load Balancing:** Round-robin, weighted, least connections algorithms
- **Caching:** In-memory, distributed caching strategies
- **Redis/Memcached:** Implementing high-performance caching
- **CDN:** Content delivery and edge computing concepts

#### **Week 11-12: Data Distribution**

- **Sharding:** Horizontal partitioning strategies
- **Consistent Hashing:** Distributed hash tables and ring topology
- **Replication:** Master-slave, master-master patterns
- **Leader-Follower Replication:** Consensus and consistency

#### **Practical Exercises:**

- Set up NGINX load balancer with multiple backend servers
- Implement Redis cluster with sharding
- Build consistent hashing algorithm
- Configure PostgreSQL streaming replication

### **2.2 Distributed Systems Theory**

#### **Week 13-14: Theoretical Foundations**

- **CAP Theorem:** Consistency, Availability, Partition tolerance trade-offs
- **Consistency Models:** Strong, eventual, causal consistency
- **Eventual Consistency:** Vector clocks, conflict resolution
- **Distributed Transactions:** Two-phase commit, Saga patterns

## Week 15-16: Advanced Patterns

- **Leader Election:** Raft, PBFT consensus algorithms
- **Fault Tolerance:** Byzantine fault tolerance, failure detection
- **Partitioning:** Network partitions and split-brain scenarios

### Practical Exercises:

- Implement Raft consensus algorithm (simplified version)
- Build eventually consistent system with conflict resolution
- Simulate network partitions and observe system behavior

## 2.3 Architecture Patterns

### Week 17-18: Service Architecture

- **Monolith vs Microservices:** Trade-offs and migration strategies
- **Microservices:** Service boundaries, data ownership
- **Service Discovery:** Registry patterns, health checks
- **Database Scaling:** Read replicas, write scaling, CQRS

### Practical Exercises:

- Decompose monolithic application into microservices
- Set up service discovery with Consul or etcd
- Implement CQRS pattern for read/write separation

---

## Phase 3: Communication & Integration (Months 8-10)

### 3.1 API Design & Management

#### Week 19-20: API Fundamentals

- **REST vs gRPC:** HTTP vs binary protocol trade-offs
- **API Versioning:** Backward compatibility strategies
- **Protocol Buffers:** Schema evolution and serialization
- **Serialization:** JSON, Avro, MessagePack performance comparison

#### Week 21-22: API Infrastructure

- **API Gateway:** Routing, transformation, aggregation
- **Rate Limiting:** Token bucket, sliding window algorithms
- **API Rate Limits:** Per-user, per-endpoint strategies
- **Throttling:** Graceful degradation under load

## Week 23: Security & Standards

- **JWT:** Stateless authentication and authorization
- **OAuth:** Delegation and third-party access
- **CORS:** Cross-origin resource sharing policies
- **API Security:** Input validation, SQL injection prevention

### Practical Exercises:

- Build REST API with OpenAPI specification
- Implement gRPC service with Protocol Buffers
- Set up Kong or AWS API Gateway
- Implement JWT authentication with refresh tokens

## 3.2 Asynchronous Communication

### Week 24: Messaging Patterns

- **Message Queues:** Point-to-point vs publish-subscribe
- **Asynchronous Processing:** Event-driven architectures
- **Dead Letter Queue:** Error handling and poison messages
- **Fan-out/Fan-in:** Scatter-gather communication patterns

### Week 25-26: Advanced Communication

- **WebSockets:** Real-time bidirectional communication
- **Long Polling:** Alternative to WebSockets for real-time updates
- **Service Mesh:** Istio, Linkerd for service communication
- **Queueing:** RabbitMQ, Apache Kafka implementation patterns

### Practical Exercises:

- Build chat application with WebSockets
  - Implement event sourcing with Apache Kafka
  - Set up Istio service mesh in Kubernetes
  - Create dead letter queue handling system
-

## Phase 4: Production & Reliability (Months 11-12)

### 4.1 Reliability Engineering

#### Week 27-28: Reliability Patterns

- **Circuit Breaker:** Preventing cascade failures
- **Retry Patterns:** Exponential backoff, jitter
- **Idempotency:** Ensuring safe retries
- **Graceful Degradation:** Maintaining core functionality under stress

#### Week 29-30: Failure Management

- **Failover:** Automatic and manual failover strategies
- **Health Checks:** Liveness and readiness probes
- **Heartbeats:** Detecting and handling node failures
- **Retry Logic:** Smart retry with circuit breaker integration

#### Practical Exercises:

- Implement circuit breaker pattern
- Build health check endpoints
- Create retry mechanism with exponential backoff
- Test failover scenarios

### 4.2 Observability & Operations

#### Week 31-32: Monitoring & Metrics

- **Metrics:** Application and infrastructure monitoring
- **Logging:** Structured logging, log aggregation
- **Distributed Tracing:** Request flow across services
- **Monitoring:** Prometheus, Grafana, ELK stack

#### Week 33: Performance & Scaling

- **Load Testing:** JMeter, K6 for performance validation
- **Autoscaling:** Horizontal and vertical scaling triggers
- **SLO/SLI/SLA:** Service level objectives and indicators
- **Error Budgets:** Balancing reliability and feature velocity

#### Practical Exercises:

- Set up monitoring stack (Prometheus + Grafana)
- Implement distributed tracing with Jaeger
- Create comprehensive load testing suite
- Define SLOs and error budgets for services

## 4.3 Deployment & Incident Management

### Week 34: Deployment Strategies

- **Blue-Green Deployment:** Zero-downtime deployments
- **Canary Deployments:** Gradual rollout strategies
- **Rollbacks:** Quick recovery from failed deployments
- **Chaos Engineering:** Proactive failure testing

### Week 35-36: Operations

- **Incident Response:** Runbooks and escalation procedures
- **Alerting:** Alert fatigue prevention and smart alerting
- **Chaos Engineering:** Netflix Chaos Monkey concepts

### Practical Exercises:

- Implement blue-green deployment pipeline
  - Set up canary deployment with automatic rollback
  - Create incident response runbooks
  - Practice chaos engineering scenarios
- 

## Practical Project Recommendations

### Project 1: Distributed E-commerce Platform (Months 4-6)

Build a microservices-based e-commerce system including:

- User service with authentication
- Product catalog with search
- Order processing with inventory management
- Payment service with external integration
- Notification service

**Technologies:** Spring Boot/Node.js, PostgreSQL, Redis, RabbitMQ, Docker

### Project 2: Real-time Analytics Platform (Months 8-10)

Create a system for processing streaming data:

- Data ingestion service
- Stream processing with Apache Kafka
- Real-time dashboards
- Historical data analysis
- Alerting system

**Technologies:** Apache Kafka, Apache Spark, InfluxDB, Grafana, Kubernetes

### Project 3: Multi-tenant SaaS Platform (Months 11-12)

Design a scalable multi-tenant application:

- Tenant isolation strategies
- Shared vs dedicated resources
- Billing and metering
- Performance monitoring per tenant
- Automated scaling

**Technologies:** AWS/GCP, Kubernetes, Istio, Prometheus, Terraform

---

## Learning Resources

### Books (Priority Reading)

1. **"Designing Data-Intensive Applications" by Martin Kleppmann** - Fundamental concepts
2. **"Building Microservices" by Sam Newman** - Service architecture
3. **"Site Reliability Engineering" by Google** - Production operations
4. **"Release It!" by Michael Nygard** - Production-ready software

### Online Resources

- **AWS Architecture Center:** Real-world architecture patterns
- **High Scalability:** Case studies from major tech companies
- **Papers We Love:** Academic papers on distributed systems
- **Kubernetes Documentation:** Container orchestration

### Tools to Master

- **Containerization:** Docker, Kubernetes
  - **Infrastructure as Code:** Terraform, CloudFormation
  - **CI/CD:** Jenkins, GitLab CI, GitHub Actions
  - **Monitoring:** Prometheus, Grafana, ELK Stack
  - **Service Mesh:** Istio, Linkerd
- 

## Career Progression Milestones

### Junior/Mid-level (Months 1-6)

- ☒ Understand database fundamentals and scaling basics
- ☒ Can design simple distributed systems
- ☒ Familiar with caching and load balancing
- ☒ Basic understanding of microservices

### Senior Level (Months 7-10)

- ☒ Design complex distributed systems
- ☒ Make informed trade-off decisions (CAP theorem)
- ☒ Implement reliable communication patterns
- ☒ Design APIs and integration strategies

### Lead/Principal Level (Months 11-12+)

- ☒ System-wide architecture decisions
  - ☒ Production reliability and incident management
  - ☒ Performance optimization and capacity planning
  - ☒ Technical leadership and mentoring
- 

## Assessment & Practice

### Monthly Self-Assessment Questions

1. Can I explain the trade-offs between consistency and availability?
2. How would I design a system to handle 1M requests per second?
3. What monitoring strategy would I implement for microservices?
4. How do I ensure data consistency across distributed services?

### Hands-on Practice Schedule



- **Week 1-2:** Theory and reading
- **Week 3:** Hands-on implementation
- **Week 4:** Project work and reflection

## **Interview Preparation Topics**

- System design scenarios (URL shortener, chat system, etc.)
- Trade-off discussions (SQL vs NoSQL, sync vs async)
- Failure scenarios and recovery strategies
- Performance optimization techniques

This guide provides a structured path to mastering technical architecture. Focus on building practical experience alongside theoretical knowledge, and don't hesitate to dive deeper into areas that align with your specific career interests.